

Machine Learning

Machine Learning Resources

[Online Resources \(/machine-learning-resources/\)](/machine-learning-resources/)

[Journal Library \(/machine-learning/journals-library/\)](/machine-learning/journals-library/)

[Datasets for Machine Learning \(/machine-learning/datasets\)](/machine-learning/datasets)

Machine Learning and Econometrics

[Resources \(/machine-learning/econometrics/resources\)](/machine-learning/econometrics/resources)

Supervised Learning Theory

[Overview \(/machine-learning/\)](/machine-learning/)

[One Variable Linear Regression \(/one-variable-linear-regression/\)](/one-variable-linear-regression/)

[Linear Algebra \(/linear-algebra-machine-learning/\)](/linear-algebra-machine-learning/)

[Multiple Variable Linear Regression \(/multi-variable-linear-regression/\)](/multi-variable-linear-regression/)

[Logistic Regression \(/logistic-regression/\)](/logistic-regression/)

[Neural Networks \(Representation\) \(/neural-networks-representation/\)](/neural-networks-representation/)

[Neural Networks \(Learning\) \(/neural-networks-learning/\)](/neural-networks-learning/)

[Applying Machine Learning \(/applying-machine-learning/\)](/applying-machine-learning/)

[Machine Learning Systems Design \(/machine-learning-systems-design/\)](/machine-learning-systems-design/)

[Support Vector Machines \(/machine-learning-svms-support-vector-machines/\)](/machine-learning-svms-support-vector-machines/)

Unsupervised Learning Theory

[Unsupervised Learning \(/machine-learning-unsupervised-learning/\)](/machine-learning-unsupervised-learning/)

[Dimensionality Reduction \(/machine-learning-dimensionality-reduction/\)](/machine-learning-dimensionality-reduction/)

[Anomaly Detection \(/machine-learning-anomaly-detection/\)](/machine-learning-anomaly-detection/)

[Recommender Systems \(/machine-learning-recommender-systems/\)](/machine-learning-recommender-systems/)

[Large Scale Machine Learning \(/machine-learning-large-scale/\)](/machine-learning-large-scale/)

[Photo OCR \(/machine-learning-photo-ocr/\)](/machine-learning-photo-ocr/)

Reinforcement Learning Theory

[Markov Decision Processes \(/machine-learning-markov-decision-processes/\)](/machine-learning-markov-decision-processes/)

[Reinforcement Learning \(/machine-learning-reinforcement-learning/\)](/machine-learning-reinforcement-learning/)

[Game Theory \(/machine-learning-game-theory/\)](/machine-learning-game-theory/)

Deep Learning Theory

[Deep Learning Terms \(/machine-learning/deep-learning/terms/\)](/machine-learning/deep-learning/terms/)

[Deep Learning Intro \(/machine-learning/deep-learning/intro/\)](/machine-learning/deep-learning/intro/)

[Deep Neural Networks Intro \(/machine-learning/deep-learning/neural-nets/\)](/machine-learning/deep-learning/neural-nets/)

[Deep Convolutional Networks Intro \(/machine-learning/deep-learning/convnets/\)](/machine-learning/deep-learning/convnets/)

Deep Learning with TensorFlow

[Exploring NotMNIST \(/machine-learning/deep-learning/tensorflow/notmnist/\)](/machine-learning/deep-learning/tensorflow/notmnist/)

[Deep Neural Networks \(/machine-learning/deep-learning/tensorflow/deep-neural-nets/\)](/machine-learning/deep-learning/tensorflow/deep-neural-nets/)

[Regularization \(/machine-learning/deep-learning/tensorflow/regularization/\)](/machine-learning/deep-learning/tensorflow/regularization/)

[Deep Convolutional Networks \(/machine-learning/deep-learning/tensorflow/convnets/\)](/machine-learning/deep-learning/tensorflow/convnets/)

Machine Learning with Scikit-Learn

[Introduction to Machine Learning \(/machine-learning-intro-easy/\)](/machine-learning-intro-easy/)

[IPython Introduction \(/ipython-introduction/\)](/ipython-introduction/)

[Iris Dataset \(/machine-learning-iris-dataset/\)](/machine-learning-iris-dataset/)

[Linear Regression Model \(/machine-learning-linear-regression/\)](/machine-learning-linear-regression/)

[Linear Regression Model Evaluation \(/machine-learning-evaluate-linear-regression-model/\)](/machine-learning-evaluate-linear-regression-model/)

[Polynomial Regression \(/machine-learning-polynomial-regression/\)](/machine-learning-polynomial-regression/)

[Vectorization, Multinomial Naive Bayes Classifier and Evaluation \(/machine-learning-multinomial-naive-bayes-vectorization/\)](/machine-learning-multinomial-naive-bayes-vectorization/)

[Gaussian Naive Bayes \(/machine-learning-gaussian-naive-bayes/\)](/machine-learning-gaussian-naive-bayes/)

[K-nearest Neighbors \(KNN\) Classification Model \(/machine-learning-k-nearest-neighbors-knn/\)](/machine-learning-k-nearest-neighbors-knn/)

[Ensemble Learning and Adaboost \(/machine-learning-ensemble-of-learners-adaboost/\)](/machine-learning-ensemble-of-learners-adaboost/)

[Decision Trees \(/machine-learning-decision-trees/\)](/machine-learning-decision-trees/)

[Support Vector Machines \(/machine-learning-svms/\)](/machine-learning-svms/)

[Clustering with KMeans \(/machine-learning-clustering-kmeans/\)](/machine-learning-clustering-kmeans/)

[Dimensionality Reduction and Feature Transformation \(/machine-learning-dimensionality-reduction-feature-transform/\)](/machine-learning-dimensionality-reduction-feature-transform/)

[Feature Engineering and Scaling \(/machine-learning-feature_engineering_scaling/\)](/machine-learning-feature_engineering_scaling/)

[Cross-Validation for Parameter Tuning, Model Selection, and Feature Selection \(/machine-learning-cross-validation/\)](/machine-learning-cross-validation/)

[Efficiently Searching Optimal Tuning Parameters \(/machine-learning-efficiently-search-tuning-param/\)](/machine-learning-efficiently-search-tuning-param/)

[Evaluating a Classification Model \(/machine-learning-evaluate-classification-model/\)](/machine-learning-evaluate-classification-model/)

[One Hot Encoding \(/machinelearning-one-hot-encoding/\)](/machinelearning-one-hot-encoding/)

[F1 Score \(/machinelearning-f1-score/\)](/machinelearning-f1-score/)

[Learning Curve \(/machinelearning-learning-curve/\)](/machinelearning-learning-curve/)

Machine Learning Projects

[Titanic Survival Data Exploration \(/machine-learning-project-titanic-survival/\)](/machine-learning-project-titanic-survival/)

[Boston House Prices Prediction and Evaluation \(Model Evaluation and Prediction\) \(/machine-learning-project-boston-home-prices/\)](/machine-learning-project-boston-home-prices/)

[Building a Student Intervention System \(Supervised Learning\) \(/machine-learning-project-student-intervention/\)](/machine-learning-project-student-intervention/)

[Identifying Customer Segments \(Unsupervised Learning\) \(/machine-learning-project-customer-segments/\)](/machine-learning-project-customer-segments/)

[Training a Smart Cab \(Reinforcement Learning\) \(/machine-learning-proj-smart-cab/\)](/machine-learning-proj-smart-cab/)

Evaluating a Classification Model

Summary: ROC, AUC, confusion matrix, and metrics

Topics

1. Review of model evaluation
2. Model evaluation procedures
3. Model evaluation metrics
4. Classification accuracy
5. Confusion matrix
6. Metrics computed from a confusion matrix
7. Adjusting the classification threshold
8. Receiver Operating Characteristic (ROC) Curves
9. Area Under the Curve (AUC)
10. Confusion Matrix Resources
11. ROC and AUC Resources
12. Other Resources

This tutorial is derived from Data School's Machine Learning with scikit-learn tutorial. I added my own notes so anyone, including myself, can refer to this tutorial without watching the videos.

1. Review of model evaluation

- Need a way to choose between models: different model types, tuning parameters, and features
- Use a **model evaluation procedure** to estimate how well a model will generalize to out-of-sample data
- Requires a **model evaluation metric** to quantify the model performance

2. Model evaluation procedures

1. Training and testing on the same data

- Rewards overly complex models that "overfit" the training data and won't necessarily generalize

2. Train/test split

- Split the dataset into two pieces, so that the model can be trained and tested on different data
- Better estimate of out-of-sample performance, but still a "high variance" estimate
- Useful due to its speed, simplicity, and flexibility

3. K-fold cross-validation

- Systematically create "K" train/test splits and average the results together
- Even better estimate of out-of-sample performance
- Runs "K" times slower than train/test split

3. Model evaluation metrics

- **Regression problems:** Mean Absolute Error, Mean Squared Error, Root Mean Squared Error
- **Classification problems:** Classification accuracy
 - There are many more metrics, and we will discuss them today

4. Classification accuracy

Pima Indian Diabetes dataset (<https://archive.ics.uci.edu/ml/datasets/Pima+Indians+Diabetes>) from the UCI Machine Learning Repository

```
In [1]: # read the data into a Pandas DataFrame
import pandas as pd

url = 'https://archive.ics.uci.edu/ml/machine-learning-databases/pima-indians-diabetes/pima-indians-diabetes.data'

col_names = ['pregnant', 'glucose', 'bp', 'skin', 'insulin', 'bmi', 'pedigree', 'age', 'label']

pima = pd.read_csv(url, header=None, names=col_names)
```

```
In [2]: # print the first 5 rows of data from the dataframe
pima.head()
```

```
Out[2]:
```

	pregnant	glucose	bp	skin	insulin	bmi	pedigree	age	label
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

- label
 - 1: diabetes
 - 0: no diabetes
- pregnant
 - number of times pregnant

Question: Can we predict the diabetes status of a patient given their health measurements?

```
In [3]: # define X and y
feature_cols = ['pregnant', 'insulin', 'bmi', 'age']

# X is a matrix, hence we use [] to access the features we want in feature_cols
X = pima[feature_cols]

# y is a vector, hence we use dot to access 'label'
y = pima.label
```

```
In [4]: # split X and y into training and testing sets
from sklearn.cross_validation import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

```
In [5]: # train a logistic regression model on the training set
        from sklearn.linear_model import LogisticRegression

        # instantiate model
        logreg = LogisticRegression()

        # fit model
        logreg.fit(X_train, y_train)
```

```
Out[5]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                           intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                           penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                           verbose=0, warm_start=False)
```

```
In [6]: # make class predictions for the testing set
        y_pred_class = logreg.predict(X_test)
```

Classification accuracy: percentage of correct predictions

```
In [7]: # calculate accuracy
        from sklearn import metrics
        print(metrics.accuracy_score(y_test, y_pred_class))

0.692708333333
```

Classification accuracy is 69%

Null accuracy: accuracy that could be achieved by always predicting the most frequent class

- We must always compare with this

```
In [8]: # examine the class distribution of the testing set (using a Pandas Series method)
        y_test.value_counts()
```

```
Out[8]: 0    130
        1     62
        Name: label, dtype: int64
```

```
In [9]: # calculate the percentage of ones
        # because y_test only contains ones and zeros, we can simply calculate the mean
        = percentage of ones
        y_test.mean()
```

```
Out[9]: 0.3229166666666667
```

32% of the

```
In [10]: # calculate the percentage of zeros
1 - y_test.mean()
```

```
Out[10]: 0.6770833333333333
```

```
In [11]: # calculate null accuracy in a single line of code
# only for binary classification problems coded as 0/1
max(y_test.mean(), 1 - y_test.mean())
```

```
Out[11]: 0.6770833333333333
```

This means that a dumb model that always predicts 0 would be right 68% of the time

- This shows how classification accuracy is not that good as it's close to a dumb model
- It's a good way to know the minimum we should achieve with our models

```
In [12]: # calculate null accuracy (for multi-class classification problems)
y_test.value_counts().head(1) / len(y_test)
```

```
Out[12]: 0    0.677083
         Name: label, dtype: float64
```

Comparing the **true** and **predicted** response values

```
In [13]: # print the first 25 true and predicted responses
print('True:', y_test.values[0:25])
print('False:', y_pred_class[0:25])
```

```
True: [1 0 0 1 0 0 1 1 0 0 1 1 0 0 0 0 1 0 0 0 1 1 0 0 0]
False: [0 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0]
```

Conclusion:

- Classification accuracy is the **easiest classification metric to understand**
- But, it does not tell you the **underlying distribution** of response values
 - We examine by calculating the null accuracy
- And, it does not tell you what **"types" of errors** your classifier is making

5. Confusion matrix

Table that describes the performance of a classification model

```
In [14]: # IMPORTANT: first argument is true values, second argument is predicted values
# this produces a 2x2 numpy array (matrix)
print(metrics.confusion_matrix(y_test, y_pred_class))

[[118  12]
 [ 47  15]]
```

n=192	Predicted: 0	Predicted: 1
Actual: 0	118	12
Actual: 1	47	15

- Every observation in the testing set is represented in **exactly one box**
- It's a 2x2 matrix because there are **2 response classes**
- The format shown here is **not** universal
 - Take attention to the format when interpreting a confusion matrix

Basic terminology

- **True Positives (TP):** we *correctly* predicted that they *do* have diabetes
 - 15
 - **True Negatives (TN):** we *correctly* predicted that they *don't* have diabetes
 - 118
 - **False Positives (FP):** we *incorrectly* predicted that they *do* have diabetes (a "Type I error")
 - 12
 - Falsely predict positive
 - Type I error
 - **False Negatives (FN):** we *incorrectly* predicted that they *don't* have diabetes (a "Type II error")
 - 47
 - Falsely predict negative
 - Type II error
-
- 0: negative class
 - 1: positive class

In [15]: *# print the first 25 true and predicted responses*

```
print('True', y_test.values[0:25])
print('Pred', y_pred_class[0:25])
```

```
True [1 0 0 1 0 0 1 1 0 0 1 1 0 0 0 0 1 0 0 0 1 1 0 0 0]
Pred [0 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0]
```

In [16]: *# save confusion matrix and slice into four pieces*

```
confusion = metrics.confusion_matrix(y_test, y_pred_class)
print(confusion)
#[row, column]
TP = confusion[1, 1]
TN = confusion[0, 0]
FP = confusion[0, 1]
FN = confusion[1, 0]
```

```
[[118 12]
 [ 47 15]]
```

		Predicted: 0	Predicted: 1	
Actual: 0	n=192	TN = 118	FP = 12	130
Actual: 1		FN = 47	TP = 15	62
		165	27	

6. Metrics computed from a confusion matrix

Classification Accuracy: Overall, how often is the classifier correct?

```
In [17]: # use float to perform true division, not integer division
print((TP + TN) / float(TP + TN + FP + FN))
print(metrics.accuracy_score(y_test, y_pred_class))

0.692708333333
0.692708333333
```

Classification Error: Overall, how often is the classifier incorrect?

- Also known as "Misclassification Rate"

```
In [18]: classification_error = (FP + FN) / float(TP + TN + FP + FN)

print(classification_error)
print(1 - metrics.accuracy_score(y_test, y_pred_class))

0.307291666667
0.307291666667
```

Sensitivity: When the actual value is positive, how often is the prediction correct?

- Something we want to maximize
- How "sensitive" is the classifier to detecting positive instances?
- Also known as "True Positive Rate" or "Recall"
- $TP / \text{all positive}$
 - $\text{all positive} = TP + FN$

```
In [19]: sensitivity = TP / float(FN + TP)

print(sensitivity)
print(metrics.recall_score(y_test, y_pred_class))

0.241935483871
0.241935483871
```

Specificity: When the actual value is negative, how often is the prediction correct?

- Something we want to maximize
- How "specific" (or "selective") is the classifier in predicting positive instances?
- $TN / \text{all negative}$
 - $\text{all negative} = TN + FP$

```
In [20]: specificity = TN / (TN + FP)

print(specificity)

0.907692307692
```

Our classifier

- Highly specific
- Not sensitive

False Positive Rate: When the actual value is negative, how often is the prediction incorrect?

```
In [21]: false_positive_rate = FP / float(TN + FP)

print(false_positive_rate)
print(1 - specificity)

0.0923076923077
0.0923076923077
```

Precision: When a positive value is predicted, how often is the prediction correct?

- How "precise" is the classifier when predicting positive instances?

```
In [22]: precision = TP / float(TP + FP)

print(precision)
print(metrics.precision_score(y_test, y_pred_class))

0.555555555556
0.555555555556
```

Many other metrics can be computed: F1 score, Matthews correlation coefficient, etc.

Conclusion:

- Confusion matrix gives you a **more complete picture** of how your classifier is performing
- Also allows you to compute various **classification metrics**, and these metrics can guide your model selection

Which metrics should you focus on?

- Choice of metric depends on your **business objective**
 - Identify if FP or FN is more important to reduce
 - Choose metric with relevant variable (FP or FN in the equation)
- **Spam filter** (positive class is "spam"):
 - Optimize for **precision or specificity**
 - precision
 - false positive as variable
 - specificity
 - false positive as variable
 - Because false negatives (spam goes to the inbox) are more acceptable than false positives (non-spam is caught by the spam filter)
- **Fraudulent transaction detector** (positive class is "fraud"):
 - Optimize for **sensitivity**
 - FN as a variable
 - Because false positives (normal transactions that are flagged as possible fraud) are more acceptable than false negatives (fraudulent transactions that are not detected)

7. Adjusting the classification threshold

```
In [23]: # print the first 10 predicted responses
# 1D array (vector) of binary values (0, 1)
logreg.predict(X_test)[0:10]
```

```
Out[23]: array([0, 0, 0, 0, 0, 0, 0, 1, 0, 1])
```

```
In [24]: # print the first 10 predicted probabilities of class membership
logreg.predict_proba(X_test)[0:10]
```

```
Out[24]: array([[ 0.63247571,  0.36752429],
 [ 0.71643656,  0.28356344],
 [ 0.71104114,  0.28895886],
 [ 0.5858938 ,  0.4141062 ],
 [ 0.84103973,  0.15896027],
 [ 0.82934844,  0.17065156],
 [ 0.50110974,  0.49889026],
 [ 0.48658459,  0.51341541],
 [ 0.72321388,  0.27678612],
 [ 0.32810562,  0.67189438]])
```

- Row: observation
 - Each row, numbers sum to 1
- Column: class
 - 2 response classes there 2 columns
 - column 0: predicted probability that each observation is a member of class 0
 - column 1: predicted probability that each observation is a member of class 1
- Importance of predicted probabilities
 - We can rank observations by probability of diabetes
 - Prioritize contacting those with a higher probability
- predict_proba process
 1. Predicts the probabilities
 2. Choose the class with the highest probability
- There is a 0.5 classification threshold
 - Class 1 is predicted if probability > 0.5
 - Class 0 is predicted if probability < 0.5

```
In [25]: # print the first 10 predicted probabilities for class 1
logreg.predict_proba(X_test)[0:10, 1]
```

```
Out[25]: array([ 0.36752429,  0.28356344,  0.28895886,  0.4141062 ,  0.15896027,
 0.17065156,  0.49889026,  0.51341541,  0.27678612,  0.67189438])
```

```
In [26]: # store the predicted probabilities for class 1
y_pred_prob = logreg.predict_proba(X_test)[: , 1]
```

```
In [57]: # allow plots to appear in the notebook
%matplotlib inline
import matplotlib.pyplot as plt


# adjust the font size
plt.rcParams['font.size'] = 12
```

```
In [58]: # histogram of predicted probabilities

# 8 bins
plt.hist(y_pred_prob, bins=8)

# x-axis limit from 0 to 1
plt.xlim(0,1)
plt.title('Histogram of predicted probabilities')
plt.xlabel('Predicted probability of diabetes')
plt.ylabel('Frequency')
```

```
Out[58]: <matplotlib.text.Text at 0x11c2b1128>
```



- We can see from the third bar
 - About 45% of observations have probability from 0.2 to 0.3
 - Small number of observations with probability > 0.5
 - This is below the threshold of 0.5
 - Most would be predicted "no diabetes" in this case
- Solution
 - **Decrease the threshold** for predicting diabetes
 - **Increase the sensitivity** of the classifier
 - This would increase the number of TP
 - More sensitive to positive instances
 - Example of metal detector
 - Threshold set to set off alarm for large object but not tiny objects
 - YES: metal, NO: no metal
 - We lower the threshold amount of metal to set it off
 - It is now more sensitive to metal
 - It will then predict YES more often

```
In [29]: # predict diabetes if the predicted probability is greater than 0.3
from sklearn.preprocessing import binarize
# it will return 1 for all values above 0.3 and 0 otherwise
# results are 2D so we slice out the first column
y_pred_class = binarize(y_pred_prob, 0.3)[0]

/Users/ritchieng/anaconda3/envs/py3k/lib/python3.5/site-packages/sklearn/utils/validation.py:386: DeprecationWarning: Passing 1d arrays as data is deprecated in 0.17 and will raise ValueError in 0.19. Reshape your data either using X.reshape(-1, 1) if your data has a single feature or X.reshape(1, -1) if it contains a single sample.
  DeprecationWarning)
```

```
In [30]: # print the first 10 predicted probabilities
y_pred_prob[0:10]
```

```
Out[30]: array([ 0.36752429,  0.28356344,  0.28895886,  0.4141062 ,  0.15896027,
                0.17065156,  0.49889026,  0.51341541,  0.27678612,  0.67189438])
```

```
In [40]: # print the first 10 predicted classes with the lower threshold
y_pred_class[0:10]
```

```
Out[40]: array([ 1.,  0.,  0.,  1.,  0.,  0.,  1.,  1.,  0.,  1.])
```



```
In [41]: # previous confusion matrix (default threshold of 0.5)
print(confusion)

[[118  12]
 [ 47  15]]
```

```
In [42]: # new confusion matrix (threshold of 0.3)
print(metrics.confusion_matrix(y_test, y_pred_class))

[[80 50]
 [16 46]]
```

- The row totals are the same
- The rows represent actual response values
 - 130 values top row
 - 62 values bottom row
- Observations from the left column moving to the right column because we will have more TP and FP

```
In [46]: # sensitivity has increased (used to be 0.24)
print (46 / float(46 + 16))

0.7419354838709677
```

```
In [44]: # specificity has decreased (used to be 0.91)
print(80 / float(80 + 50))

0.6153846153846154
```

Conclusion:

- **Threshold of 0.5** is used by default (for binary problems) to convert predicted probabilities into class predictions
- Threshold can be **adjusted** to increase sensitivity or specificity
- Sensitivity and specificity have an **inverse relationship**
 - Increasing one would always decrease the other
- Adjusting the threshold should be one of the last step you do in the model-building process
 - The most important steps are
 - Building the models
 - Selecting the best model

8. Receiver Operating Characteristic (ROC) Curves

Question: Wouldn't it be nice if we could see how sensitivity and specificity are affected by various thresholds, without actually changing the threshold?


Answer: Plot the ROC curve.

- Receiver Operating Characteristic (ROC)

```
In [59]: # IMPORTANT: first argument is true values, second argument is predicted probabilities

# we pass y_test and y_pred_prob
# we do not use y_pred_class, because it will give incorrect results without generating an error
# roc_curve returns 3 objects fpr, tpr, thresholds
# fpr: false positive rate
# tpr: true positive rate
fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred_prob)

plt.plot(fpr, tpr)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.rcParams['font.size'] = 12
plt.title('ROC curve for diabetes classifier')
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.grid(True)
```



- ROC curve can help you to **choose a threshold** that balances sensitivity and specificity in a way that makes sense for your particular context
- You can't actually **see the thresholds** used to generate the curve on the ROC curve itself

```
In [63]: # define a function that accepts a threshold and prints sensitivity and specificity
def evaluate_threshold(threshold):
    print('Sensitivity:', tpr[thresholds > threshold][-1])
    print('Specificity:', 1 - fpr[thresholds > threshold][-1])
```

```
In [64]: evaluate_threshold(0.5)

Sensitivity: 0.241935483871
Specificity: 0.907692307692
```

```
In [65]: evaluate_threshold(0.3)

Sensitivity: 0.725806451613
Specificity: 0.615384615385
```

9. AUC

AUC is the **percentage** of the ROC plot that is **underneath the curve**:

```
In [67]: # IMPORTANT: first argument is true values, second argument is predicted probabilities
print(metrics.roc_auc_score(y_test, y_pred_prob))

0.724565756824
```

- AUC is useful as a **single number summary** of classifier performance
- Higher value = better classifier
- If you randomly chose one positive and one negative observation, AUC represents the likelihood that your classifier will assign a **higher predicted probability** to the positive observation
- AUC is useful even when there is **high class imbalance** (unlike classification accuracy)
 - Fraud case
 - Null accuracy almost 99%
 - AUC is useful here

```
In [68]: # calculate cross-validated AUC
from sklearn.cross_validation import cross_val_score
cross_val_score(logreg, X, y, cv=10, scoring='roc_auc').mean()
```

```
Out[68]: 0.73782336182336183
```

Use both of these whenever possible

1. Confusion matrix advantages:

- Allows you to calculate a **variety of metrics**
- Useful for **multi-class problems** (more than two response classes)

2. ROC/AUC advantages:

- Does not require you to **set a classification threshold**
- Still useful when there is **high class imbalance**

10. Confusion Matrix Resources

- Blog post: [Simple guide to confusion matrix terminology](http://www.dataschool.io/simple-guide-to-confusion-matrix-terminology/) (<http://www.dataschool.io/simple-guide-to-confusion-matrix-terminology/>) by me
- Videos: [Intuitive sensitivity and specificity](https://www.youtube.com/watch?v=U4_3fditnWg) (https://www.youtube.com/watch?v=U4_3fditnWg) (9 minutes) and [The tradeoff between sensitivity and specificity](https://www.youtube.com/watch?v=vtYDyGGeQyo) (<https://www.youtube.com/watch?v=vtYDyGGeQyo>) (13 minutes) by Rahul Patwari
- Notebook: [How to calculate "expected value"](https://github.com/podopie/DAT18NYC/blob/master/classes/13-expected_value_cost_benefit_analysis.ipynb) (https://github.com/podopie/DAT18NYC/blob/master/classes/13-expected_value_cost_benefit_analysis.ipynb) from a confusion matrix by treating it as a cost-benefit matrix (by Ed Podojil)
- Graphic: How [classification threshold](https://media.amazonwebservices.com/blog/2015/ml_adjust_model_1.png) (https://media.amazonwebservices.com/blog/2015/ml_adjust_model_1.png) affects different evaluation metrics (from a [blog post](https://aws.amazon.com/blogs/aws/amazon-machine-learning-make-data-driven-decisions-at-scale/) (<https://aws.amazon.com/blogs/aws/amazon-machine-learning-make-data-driven-decisions-at-scale/>) about Amazon Machine Learning)

11. ROC and AUC Resources

- Lesson notes: [ROC Curves](http://ebp.uga.edu/courses/Chapter%204%20-%20Diagnosis%20I/8%20-%20ROC%20curves.html) (<http://ebp.uga.edu/courses/Chapter%204%20-%20Diagnosis%20I/8%20-%20ROC%20curves.html>) (from the University of Georgia)
- Video: [ROC Curves and Area Under the Curve](https://www.youtube.com/watch?v=OAl6eAyP-yo) (<https://www.youtube.com/watch?v=OAl6eAyP-yo>) (14 minutes) by me, including [transcript and screenshots](http://www.dataschool.io/roc-curves-and-auc-explained/) (<http://www.dataschool.io/roc-curves-and-auc-explained/>) and a [visualization](http://www.navan.name/roc/) (<http://www.navan.name/roc/>)
- Video: [ROC Curves](https://www.youtube.com/watch?v=21Igj5Pr6u4) (<https://www.youtube.com/watch?v=21Igj5Pr6u4>) (12 minutes) by Rahul Patwari
- Paper: [An introduction to ROC analysis](http://people.inf.elte.hu/kiss/13dwhdm/roc.pdf) (<http://people.inf.elte.hu/kiss/13dwhdm/roc.pdf>) by Tom Fawcett
- Usage examples: [Comparing different feature sets](http://research.microsoft.com/pubs/205472/aisec10-leontjeva.pdf) (<http://research.microsoft.com/pubs/205472/aisec10-leontjeva.pdf>) for detecting fraudulent Skype users, and [comparing different classifiers](http://www.cse.ust.hk/nevinZhangGroup/readings/yi/Bradley_PR97.pdf) (http://www.cse.ust.hk/nevinZhangGroup/readings/yi/Bradley_PR97.pdf) on a number of popular datasets

12. Other Resources

- scikit-learn documentation: [Model evaluation](http://scikit-learn.org/stable/modules/model_evaluation.html) (http://scikit-learn.org/stable/modules/model_evaluation.html)

- Guide: [Comparing model evaluation procedures and metrics](https://github.com/justmarkham/DAT8/blob/master/other/model_evaluation_comparison.md) (https://github.com/justmarkham/DAT8/blob/master/other/model_evaluation_comparison.md) by me
- Video: [Counterfactual evaluation of machine learning models](https://www.youtube.com/watch?v=QWCSxAKR-h0) (<https://www.youtube.com/watch?v=QWCSxAKR-h0>) (45 minutes) about how Stripe evaluates its fraud detection model, including [slides](http://www.slideshare.net/MichaelManapat/counterfactual-evaluation-of-machine-learning-models) (<http://www.slideshare.net/MichaelManapat/counterfactual-evaluation-of-machine-learning-models>)

Tags:

9 Comments

Ritchie Ng | Deep Learning & Computer Vision Engineer

 Login ▾

 Recommend 2

 Tweet

 Share

Sort by Best ▾



Join the discussion...

LOG IN WITH

OR SIGN UP WITH DISQUS 



anup parmar • a year ago

Hi Ritchie,

This is excellent explanation on evaluation of models.
Thank you for writing and sharing this code. It really help me :)

25 ^ | ▾ • Reply • Share ›



saaurabh keshari • 2 months ago

Amazing work !!!

I was desperately looking for these performance matrix.

Thanks a ton Ritchie :)

1 ^ | ▾ • Reply • Share ›



Muhammad Nadeem Ferozi • 3 months ago

Thanks a million

1 ^ | ▾ • Reply • Share ›



General Useage • a year ago

AMAZING THANK YOU

1 ^ | ▾ • Reply • Share ›



Praful Hanbarde • 8 months ago

best tutorial...

each concept explain very clearly

^ | v • Reply • Share ›

Alfredo Contreras • a year ago

This is a great post! Thanks very much for sharing it.

^ | v • Reply • Share ›

Arjun Anil • a year ago

Hi Ritchie,

Thank you for sharing this.

I have a query related to multi class classification models.

Will the confusion matrix, precision and recall, ROC and AUC still be relevant when taking the case of multi class classification?

^ | v • Reply • Share ›



Ritchie Ng Moderator → Arjun Anil • a year ago

For confusion matrix, it is highly relevant for multi-class problems. For example you can see if you predict class 0 when the ground truth is class 4. And that side gap may be severe and you can look at minimizing this versus predicting class 0 and ground truth being 1. This is highly relevant for say cancer detection where the higher classes may be more severe cases.

For AUC, you need to compare by 2 classes and analyse accordingly.

^ | v • Reply • Share ›



Arjun Anil → Ritchie Ng • a year ago

Hi Ritchie,

Thanks for the reply. The concept has become more clear.

It would be helpful if you could make clarity on the below mentioned point:

- Is the confusion matrix only used for a train-test split method?

Can it be used in a situation where the whole data is subjected to a K-fold cross validation method?

^ | v • Reply • Share ›

ALSO ON RITCHIE NG | DEEP LEARNING & COMPUTER VISION ENGINEER

Convolutional Neural Networks with TensorFlow | Machine Learning, Deep

7 comments • 2 years ago



ayush kaul — I have tried problem 1. I got Kernal Restart error in jupyter notebook. I am running on CPU i3 4 GB memory

Comic Series: Life's Journey is Like Gradient

Changing data type | Machine Learning, Deep Learning, Reinforcement Learning and

1 comment • 2 years ago



Rinku Singh — Hello, I tried running the below query but it is failing with error `orders['item_price'] = orders.item_price.str.replace('$', ...`

Deep Convolutional Networks | Machine

Descent | Machine Learning, Deep Learning,

1 comment • a year ago



Marshall Keyes — Hi Ritchie, I can across your very interesting blog and thought I would reach out to you about something that we may share an interest

Learning, Deep Learning, Reinforcement

1 comment • 2 years ago



Rafael Padilla — Ritchie, just a little correction in your formula "Calculating Ouput Size": Where you read "-2P" should be "+2P". The padding increases

© 2018 Ritchie Ng. All rights reserved.

Site last updated: Oct 29, 2018

Github (<https://github.com/ritchieng>) | Linkedin (<https://www.linkedin.com/in/ritchieng>) | Facebook (<https://www.facebook.com/ritchiengz>) | Twitter (<https://twitter.com/ritchieng>) | Tech in Asia (<https://www.techinasia.com/profile/ritchieng>)